

RESEARCH ARTICLE

Thinking like a naturalist: Enhancing computer vision of citizen science images by harnessing contextual data

J. Christopher D. Terry^{1,2}  | Helen E. Roy¹  | Tom A. August¹ ¹NERC Centre for Ecology & Hydrology, Wallingford, UK²Department of Zoology, University of Oxford, Oxford, UK**Correspondence**J. Christopher D. Terry
Email: james.terry@zoo.ox.ac.uk**Funding information**

Natural Environment Research Council, Grant/Award Number: NE/L002612/1 and NE/R016429/1

Handling Editor: Res Altwegg**Abstract**

1. The accurate identification of species in images submitted by citizen scientists is currently a bottleneck for many data uses. Machine learning tools offer the potential to provide rapid, objective and scalable species identification for the benefit of many aspects of ecological science. Currently, most approaches only make use of image pixel data for classification. However, an experienced naturalist would also use a wide variety of contextual information such as the location and date of recording.
2. Here, we examine the automated identification of ladybird (Coccinellidae) records from the British Isles submitted to the UK Ladybird Survey, a volunteer-led mass participation recording scheme. Each image is associated with metadata; a date, location and recorder ID, which can be cross-referenced with other data sources to determine local weather at the time of recording, habitat types and the experience of the observer. We built multi-input neural network models that synthesize metadata and images to identify records to species level.
3. We show that machine learning models can effectively harness contextual information to improve the interpretation of images. Against an image-only baseline of 48.2%, we observe a 9.1 percentage-point improvement in top-1 accuracy with a multi-input model compared to only a 3.6% increase when using an ensemble of image and metadata models. This suggests that contextual data are being used to interpret an image, beyond just providing a prior expectation. We show that our neural network models appear to be utilizing similar pieces of evidence as human naturalists to make identifications.
4. Metadata is a key tool for human naturalists. We show it can also be harnessed by computer vision systems. Contextualization offers considerable extra information, particularly for challenging species, even within small and relatively homogeneous areas such as the British Isles. Although complex relationships between disparate sources of information can be profitably interpreted by simple neural network architectures, there is likely considerable room for further progress. Contextualizing images has the potential to lead to a step change in the accuracy of automated

This is an open access article under the terms of the Creative Commons Attribution License, which permits use, distribution and reproduction in any medium, provided the original work is properly cited.

© 2019 The Authors. *Methods in Ecology and Evolution* published by John Wiley & Sons Ltd on behalf of British Ecological Society.

identification tools, with considerable benefits for large-scale verification of submitted records.

KEYWORDS

citizen science, computer vision, convolutional neural network, ladybird, machine learning, metadata, naturalists, species identification

1 | INTRODUCTION

Large-scale and accurate biodiversity monitoring is a cornerstone of understanding ecosystems and human impacts upon them (IPBES, 2019). Recent advances in artificial intelligence have revolutionized the outlook for automated tools to provide rapid, scalable, objective and accurate species identification and enumeration (Torney et al., 2019; Wäldchen & Mäder, 2018; Weinstein, 2018; Willi et al., 2019). Improved accuracy levels could revolutionize the capacity of biodiversity monitoring (Isaac, Strien, August, Zeeuw, & Roy, 2014) and invasive non-native species surveillance programmes (August et al., 2015). Nonetheless, at present, general-purpose automated classification of animal species is still some distance from the level of accuracy obtained by human experts. Recent studies have achieved percentage classification accuracy ranging between mid-60s to high 90s depending on the difficulty of the problem (Wäldchen & Mäder, 2018; Weinstein, 2018), and their potential remains underutilized (Christin, Hervet, & Lecomte, 2019).

The large data requirements and capacity of machine learning has led to a close association with citizen science projects (Wäldchen & Mäder, 2018), where volunteers contribute scientific data (Silvertown, 2009). Citizen scientists can accurately crowd-source identification of researcher-gathered images (e.g. Snapshot Serengeti; Swanson et al., 2015), generate records to be validated

by experts (e.g. iRecord; Pocock, Roy, Preston, & Roy, 2015) or both simultaneously (e.g. iNaturalist; iNaturalist.org). However, there can be a considerable lag between record submission and human verification. If computer vision tools could generate more rapid, or even instantaneous, identifications it could assist with citizen scientist recruitment and retention. While image acquisition by researchers can be directly controlled and lead to high accuracies (Marques et al., 2018; Rzanny, Seeland, Wäldchen, & Mäder, 2017), images from citizen science projects are highly variable and pose considerable challenges for computer vision (Van Horn et al., 2017).

Most automatic species identification tools only make use of images (Weinstein, 2018). However, an experienced naturalist would utilize a wide variety of contextual information when making an identification. This is particularly the case when distinguishing 'difficult' species, where background information about the record may be essential for a confident identification. In a machine learning context, this Supporting Information about an image (metadata) can be split into two categories (Figure 1). Primary metadata is directly associated with a record such as GPS-coordinates, date of recording and the identity of the recorder. Derived (secondary) metadata is generated through cross-referencing with other sources of information to place this metadata into a more informative context (Tang, Paluri, Fei-Fei, Fergus, & Bourdev, 2015). In an ecological context, this may include weather records, maps of species distribution, climate or habitat, phenology records, recorder experience, or any other information source that could support an identification.

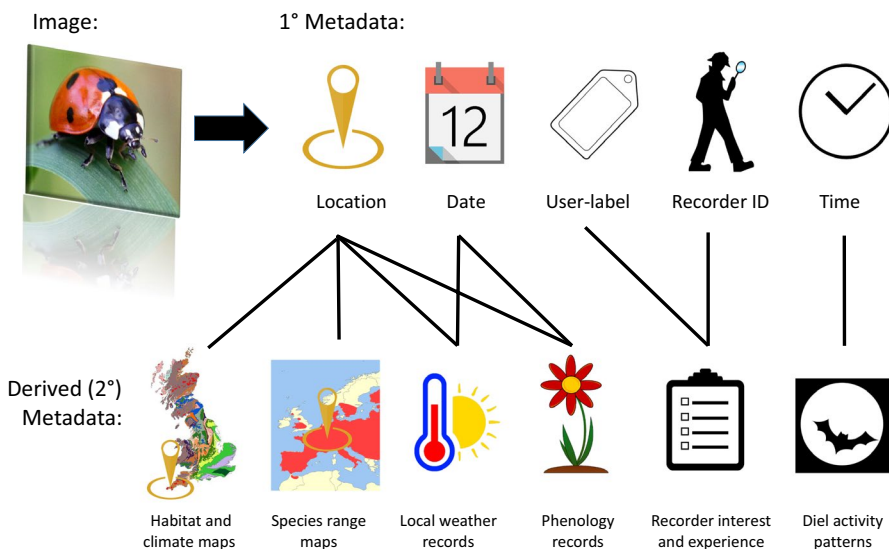


FIGURE 1 Relationships between categories of metadata. Primary metadata are basic attributes of the record directly associated with an image such as the date or location. By contrast, derived (or secondary) metadata requires cross-reference to external databases, which may include physical, ecological or social data. External sources of information may be fixed and stable (such as habitat maps) or dynamic and require updating in order to keep the model up to date (such as weather records or recorder experience)

Efforts to include contextual spatio-temporal information have largely focused on reducing the list of potential species that may be expected in a given area. iRecord (www.brc.ac.uk/irecord) partially automates this process, flagging records to expert verifiers that are labelled as being outside of the known range. Distribution priors have been shown to be effective in improving the identification of North American birds (Berg et al., 2014), images in the iNaturalist dataset (Mac Aodha, Cole, & Perona, 2019) and generating location-specific shortlists of German plants (Wittich, Seeland, Wäldchen, Rzanny, & Mäder, 2018). This approach can greatly reduce the risk of non-sensical identifications that otherwise lead to considerable scepticism over the use of automated methods (Gaston & O'Neill, 2004). Nevertheless, this 'filtering' approach does not make full use of the available data, and it has been recently shown that improvements in the identification of plankton from images can be improved by incorporating sample metadata directly into a neural network (Ellen, Graff, & Ohman, 2019). Many species vary in appearance seasonally or across their range. For example, the proportion of the melanic form of the two-spot ladybird *Adalia bipunctata* varies greatly across the UK (Creed, 1966). To an expert naturalist, metadata can do more than shorten the list of potential identifications—it can help to interpret the image itself. For example, juveniles, flowers or breeding plumage may only be observed in narrow time windows or there may be geographic variation in colour patterns. Consequently, certain features within an image (e.g. spots on a butterfly's wing) may only aid in determining a species in specific regions, or times of year. It would only be worth looking for a particular pattern when that species and life stage is active. Synthesizing and making use of such disparate sets of information is challenging for humans even when detailed data is available, and such expertise requires many years to build. By contrast, neural networks are ideally suited to drawing together diverse sources in such a way to gain the maximal amount of information.

Ladybirds (Coleoptera: Coccinellidae) are a charismatic insect family that garner substantial public interest, with large numbers of submitted records to citizen science monitoring schemes around the world (Gardiner et al., 2012). Identification of ladybirds is challenging for both human (Jouveau, Delaunay, Vignes-Lebbe, & Nattier, 2018) and artificial intelligence (Van Horn et al., 2017) because of a number of morphological features. Many species of ladybird have polymorphic elytral colour patterns, with some species seemingly mimicking others, and so are principally disambiguated by size. However, size is extremely challenging for artificial intelligence to automatically infer from a single image without standardized scales (Laina, Rupprecht, Belagiannis, Tombari, & Navab, 2016). As an example the invasive Harlequin ladybird *Harmonia axyridis* (which has been a particular focus for research, Roy et al., 2016), is a polymorphic species and can resemble a number of other species. Consequently, the Harlequin ladybird is frequently misidentified by citizen scientists (Gardiner et al., 2012) but can be distinguished on the basis of its large size and combination of other morphological features. Currently, submissions to the UK Ladybird Survey (www.ladybird-survey.org) are managed by a

small number of expert verifiers. The survey receives many tens of thousands of records every year and so the commitment required from each expert verifier is high. There is growing interest in expanding the geographic scope of the survey with the recent launch of a smartphone app for recording ladybirds across Europe (<https://european-ladybirds.brc.ac.uk/>). The UK Ladybird Survey (and associated European extension) therefore represents a real-world example of a programme where a reliable automated identification tool could help to increase the utility of citizen science to document biodiversity across the globe.

Classification tools that only use image data are not making maximal use of the information available to human experts. Here, we demonstrate methods to incorporate metadata directly within neural networks used for the classification of images of ladybirds submitted to the UK Ladybird Survey. We examine if metadata can significantly improve classification accuracy, thereby increasing their potential to assist in large-scale biodiversity monitoring, by the following:

1. Comparing the classification accuracy of classifiers incorporating metadata compared to image-only classifiers.
2. Exploring whether neural networks make use of the same pieces of metadata information that a human experts do.

2 | MATERIALS AND METHODS

2.1 | Data

Records of ladybirds (Coccinellidae) were sourced from the UK Biological Records Centre (www.brc.ac.uk). These were filtered to include only those from within the British Isles, from 2013 to 2018 inclusive, that contained an image and had been verified by an expert assessor. Records were distributed across the whole of the British Isles, although records were more frequent near more heavily populated areas (Figure S1). The date range was selected based on a notable increase in records from 2013 with the release of a mobile app (iRecord Ladybirds). Identifications of records by expert verifiers was based on uploaded images and associated information including the species determination of the original observer, location, date, associated comments and (where known) the degree of skill of the recorder.

Of the 47 species of ladybird that had been recorded at least once in the UK (Duff, 2018), only 18 species (listed in Table 1) had at least 170 usable records, which we took as our lower cut-off to ensure each species was represented by at least 120 unique training images. We judged that fewer training images would not result in accurate classification. These 18 species made up 97% of the total ladybird records during 2013–2018. Even after removing species with fewer than 170 usable records, the dataset is highly imbalanced (Table 1), with two species making up the bulk of records: seven-spot ladybird *Coccinella septempunctata* (25.8%) and the highly polymorphic Harlequin ladybird (44.5%).

TABLE 1 Average per-species top-1 accuracy across the suite of models. Citizen scientist accuracy is determined by frequency by which the label assigned by the recorder corresponds to the verified species name. Equivalent tables for top-3 accuracy and for accuracy including a prior weighting based on relative frequency are given in Tables S2 and S3. The top performing model in each row is marked with an asterisk (*)

Species	Relative frequency	Citizen scientist	Metadata only		Image only	Image and metadata	
			Primary	Derived		Combined	Ensemble
Overall		92.4	15.9	22.4	48.2	57.3*	53.7
<i>Adalia bipunctata</i>	5.3	97.3	10.9	22.5	56.4	58.9*	58.5
<i>Adalia decempunctata</i>	2.9	85.8	1.9	2.2	24.6*	22.9	23.3
<i>Anatis ocellata</i>	0.5	94.5	2.7	15.3	37.3	41.3	42.0*
<i>Aphidecta oblitterata</i>	0.7	96.5	63.2	55.3	71.6	80.0	81.6*
<i>Calvia quattuordecimguttata</i>	1.8	92.5	0.2	3.4	70.0*	55.8	69.8
<i>Chilocorus renipustulatus</i>	1.2	93.2	3.1	16.9	47.6	47.0	49.6*
<i>Coccinella septempunctata</i>	26.1	95.5	0.0	0.3	64.8*	64.2	62.9
<i>Coccinella undecimpunctata</i>	0.7	94.0	5.1	27.2	58.5	58.5	62.1*
<i>Exochomus quadripustulatus</i>	1.5	92.0	15.2	26.9	37.9	43.9*	40.0
<i>Halyzia sedecimguttata</i>	3.7	93.9	0.8	7.7	65.6	73.5*	66.1
<i>Harmonia axyridis</i>	44.1	89.6	27.9	38.6	34.7	53.6*	47.1
<i>Harmonia quadripunctata</i>	0.4	94.3	6.2	12.3	37.7	43.1*	39.2
<i>Hippodamia variegata</i>	0.6	93.2	35.4	32.0	28.0	46.9*	38.9
<i>Propylea quattuordecimpunctata</i>	4.5	94.8	28.1	22.3	58.6	62.7*	59.3
<i>Psyllobora vigintiduopunctata</i>	3.0	98.5	5.2	11.8	56.3	58.5*	58.1
<i>Scymnus interruptus</i>	0.4	98.2	93.6	76.0	88.0	89.6	90.4*
<i>Subcoccinella vigintiquattropunctata</i>	1.6	96.2	6.2	17.8	62.6	67.3*	64.5
<i>Tytthaspis sedecimpunctata</i>	1.0	91.2	7.0	21.9	43.5	51.1*	50.8

2.2 | Images

Records were manually scanned to remove the majority of images predominantly of eggs, larvae or pupae, 'contextual' images of habitat area, images including multiple species, and images that had been uploaded repeatedly. Larval and pupal images were overwhelming dominated by the highly distinctive Harlequin ladybird larvae or pupae (78%). Where a single record had multiple associated images, only the first was used. Images were centre cropped to square and then rescaled to 299×299 pixels. Example images for each species are shown in Figure 2. After all data cleaning steps, the dataset had 39,877 records in total.


2.3 | Metadata

We constructed models that made use of different subsets of the available metadata. The first (the primary metadata model) took

only three pieces of primary metadata, drawn directly from the UK Ladybird Survey dataset: longitude, latitude and date. We represented date by day-of-year, excluding year values since information on 'year' would not be transferable to future records. The second model (the derived metadata model) supplemented the primary metadata with secondary metadata: data generated with additional reference to external sources of information, namely weather records, habitat and recorder expertise. We did not use the original citizen scientist species determination in our models, since it was too powerful compared to other sources of information (correct over 92% of the time) and did not align with the goal of fully automated identification.

Temperature records were accessed from the MIDAS database (Met Office, 2012), selecting data from the 88 UK stations with fewer than 20 missing records (2013–2018). Occasional missing values were imputed with a polynomial spline. Using the closest weather station to the record, maximum daily temperature for each day in the 14 preceding days ($d-1:d-15$) and weekly average maximum

FIGURE 2 Three randomly selected images from each of the 18 ladybird species in our dataset, demonstrating the wide variety of poses, sizes and backgrounds. Images have been centre cropped to square and resized to 299×299 . Species are listed alphabetically: Left column: (a) *Adalia bipunctata*, (b) *Adalia decempunctata*, (c) *Anatis ocellata*, (d) *Aphidecta oblitterata*, (e) *Calvia quattuordecimguttata*, (f) *Chilocorus renipustulatus*, (g) *Coccinella septempunctata*, (h) *Coccinella undecimpunctata*, (i) *Exochomus quadripustulatus*. Right column: (a) *Halyzia sedecimguttata*, (b) *Harmonia axyridis*, (c) *Harmonia quadripunctata*, (d) *Hippodamia variegata*, (e) *Propylea quattuordecimpunctata*, (f) *Psyllobora vigintiduopunctata*, (g) *Scymnus interruptus*, (h) *Subcoccinella vigintiquattropunctata*, (i) *Tytthaspis sedecimpunctata*



daily temperatures for each of the 8 weeks preceding the high-resolution period ($d-16:d-71$) were accessed.

Local habitat information was derived from a 1 km resolution land cover map (Rowland et al., 2017). This provides percentages in each 1 km grid of 21 target habitat classes (e.g. 'urban', 'coniferous woodland', 'heather', etc.). Where no data was available, each habitat was assumed to be 0.

We calculated a 'recorder experience' variable as the cumulative count of records submitted by that recorder at the time of each record. Only records of ladybirds in our dataset were included in this count. Where no unique recorder ID was available, that record was assumed to be a first record.

This led to a one-dimensional metadata vector of length 47 (day-of-year, latitude, longitude, 14 daily maximum temperature records, 8 weekly average temperature records, 21 habitat frequencies and recorder experience) associated with each image.

2.4 | Machine learning model architecture

We built and fit convolutional neural network models (Goodfellow, Bengio, & Courville, 2016) in R 3.5.3 using the functional model framework of the *KERAS* package (Allaire & Chollet, 2019). We used the TensorFlow backend on a Nvidia GTX 1080 Ti GPU. R code used to train the models is available at github.com/jcdterry/LadybirdID_Public and the core model architecture code is summarized in Supporting Information. We first constructed and trained image-only and metadata-only models. Once these had separately attained maximum performance, these were then combined to form the core of a multi-input model that takes both an image and metadata as input variables. For all models, we conducted extensive hyperparameter searches to determine model architecture, extent of data-augmentation, regularization parameters, learning rates and training times.

A schematic of the model architectures is shown in Figure 3. The metadata models were built with a simple architecture of two densely connected layers and a softmax classifier layer. For the image-model, the Inception-ResNet-v2 architecture (Szegedy, Ioffe, Vanhoucke, & Alemi, 2016) was used as an initial feature extractor. This is a very deep architecture that had been pre-trained on the large imageNet dataset to extract meaningful features from a generic set of images. This transfer learning approach greatly expedites the training process and has previously achieved high accuracy in tests on the iNaturalist dataset of citizen science records (e.g. Cui, Song, Sun, Howard, & Belongie, 2018) and for the identification of insects (Martineau, Raveaux, Chatelain, Conte, & Venturini, 2018). To repurpose the model, we replaced the imageNet classification layer with new layers and trained the model on our dataset. The combined model was built by removing the classifier layers from the metadata and image models, concatenating the two outputs, and adding further layers. This fusion approach has been successfully used in the categorization of satellite data (Minetto, Pamplona Segundo, & Sarkar, 2019).

2.5 | Model training

Species records in the UK Ladybird Survey, like most biological record datasets (Van Horn et al., 2017), are highly skewed towards certain common species (Table 1). As predictive models are not perfect, such class-imbalanced data leads to critical choices about how to best assess 'accuracy'. Overall accuracy may be maximized by rarely or never assigning species to infrequent categories. A citizen scientist may prefer the maximum accuracy for the species in front of them (which is likely to be a commonly reported species). However, in an ecological science context, rare (or more precisely, rarely reported) species are often of particular interest to researchers managing citizen science projects.

The total dataset was randomly partitioned into training (70%), validation (15%) and test (15%) sets. To address the class-imbalance, we followed the approach suggested by Buda, Maki, and Mazurowski (2018) and rebalanced our training set through up-sampling and down-sampling the available records. We did this so that each species had 2,000 effective training records. To ensure a consistent batch-size of 32, we removed records of the most common species where necessary. Consequently, our underlying models did not have direct access to the information that, all else being equal, certain species are far more likely than others. This reduces the potential for the model 'cheating' during training by fixating on common species and ignoring rare species. To demonstrate the potential to improve overall accuracy by taking into account the relative frequency of each species, we tested weighted versions of each of the models. In these, the relative probability assigned to each species from each unweighted model (P_i) was scaled by the relative frequency of each of the species (F_i) in the training data as: $P_{weighted_i} \propto P_i F_i$.

To reduce overfitting, we made extensive use of image augmentation, weight regularization, batch normalization, dropout layers during training and introduced Gaussian noise on the metadata vector. Training optimization was based on a categorical cross-entropy loss function using the 'Adam' adaptive moment estimation optimizer (Kingma & Ba, 2014). During training, if validation loss had reached a plateau, learning rate was reduced automatically. Training was stopped (and the best model restored) if there had been no further improvement in validation loss over at least four epochs.

After fitting the derived metadata, image-only and combined models, a simple ensemble model taking a weighted average of the derived metadata and image-only model predictions was also constructed and tested. This could be considered equivalent to using the metadata to construct a prior expectation for the predictions of the image model:

$$P_{ensemble_i} \propto (1-\omega)P_{image_i} + \omega P_{meta_i},$$

where the weighting (ω) between the metadata and image model probabilities was determined by optimizing the ensemble model top-1 accuracy on the validation set.




FIGURE 3 Outline schematic of the difference in model architectures between the single input models that take either just metadata (a) or image (b) information, and the two multi-input models combining (c) or ensembling (d) both data sources. Dense layers are the principle component of neural networks, that fit linkages between every input and output node. All our dense layers incorporated a rectified linear unit (ReLU) nonlinear activation function. Inception-ResNet-v2 is a very deep feature extraction model incorporating many convolutional layers and originally trained to classify a diverse set of objects, that we refined by retraining on our ladybird dataset. The global max pooling stage summarizes the outputs of the image feature extractor for further computation by dense layers. Softmax layers output a vector that sums to one, which can be interpreted as probabilities of each potential category. Dropout, noise, batch normalization and other regularization features enacted only during training time are not shown here for simplicity. R code to build models using the *KERAS* R package (Allaire & Chollet, 2019) is given in Supporting Information, which also details further hyperparameters such as the size of the each layer

2.6 | Model testing and evaluation

Overall and species-level model performance was assessed in terms of top-1 (was the true ID rated most likely) and top-3 (was the true

ID amongst the three options rated most highly) accuracy. Because model accuracy will be dependent on the split of data into testing and training sets, and because model optimization is a non-deterministic process, we repeated the entire model fitting process five

times. For each repeat, assignment of images to training, validation and test sets was randomized.

2.7 | Role of metadata components

To examine the dependence of the model on each aspect of the metadata, we examined the decline in top-3 accuracy for each species when elements of metadata were randomized by reshuffling sets of values within the test set. We did this separately for the spatial coordinates, day-of-year, temperatures data, habitats data and recorder expertise.

3 | RESULTS

Across each of our training-test split realizations, combined multi-input models showed a marked and consistent improvement on both the image-only (+9.1 percentage points) and the ensemble models (+3.6 percentage points) (Figure 4). Species-level accuracies (averaged across the 5 split realizations) for each of the models are reported in Table 1. There was no correlation between the species-specific accuracy of the metadata-only model and the image-only model (Spearman's rank correlation test $\rho = 0.23$, $p = .34$). There was, however, a strong correlation at a per-species level between the fraction correctly identified by the original citizen-scientist recorder and the combined model ($\rho = 0.65$, $p = .003$). The combined model slightly increases the confidence assigned to the correct answer

compared to the image-only model, whereas the ensemble model leads to a decline in confidence (Appendix 3).

The overall accuracy of all models could be greatly improved by weighting the output probabilities by the prior expectation given the relative frequency of each species. For example, the average top-1 accuracy of the combined model rises from 57% to 69%. The model ranking in terms of overall accuracy was maintained (Table S2). However, these gains are made at the cost of very infrequently identifying rare species correctly. With a weighted model the two most commonly observed species, Harlequin and Seven-spot ladybirds, are correctly identified 90% and 89% of the time respectively. However, 12 infrequently observed species are correctly identified in less than 12% of cases.

The derived metadata model had an overall top-3 accuracy of 43.7% and was making at least some use of all the components of the metadata since randomizing each group caused a decline in accuracy. Accuracy of the metadata-only model peaked spatially away from the south-east of the British Isles and outside of summer (Figure 5). Metadata accuracy (43.7%) was most related to temperature. This is demonstrated by a 10% percentage point decrease in accuracy when temperature was removed. Where both temperature and day-of-year data was available, the temperature data appears to be used more (10% and 0.2% decreases respectively). It is not possible to determine whether this is because temperature is simply more relevant to ladybirds than date, or whether this is an artefact of the different lengths of the metadata vectors. When day-of-year was randomized in the primary metadata model, top-3 accuracy declines by 4.5 percentage points. Within temperature, the model

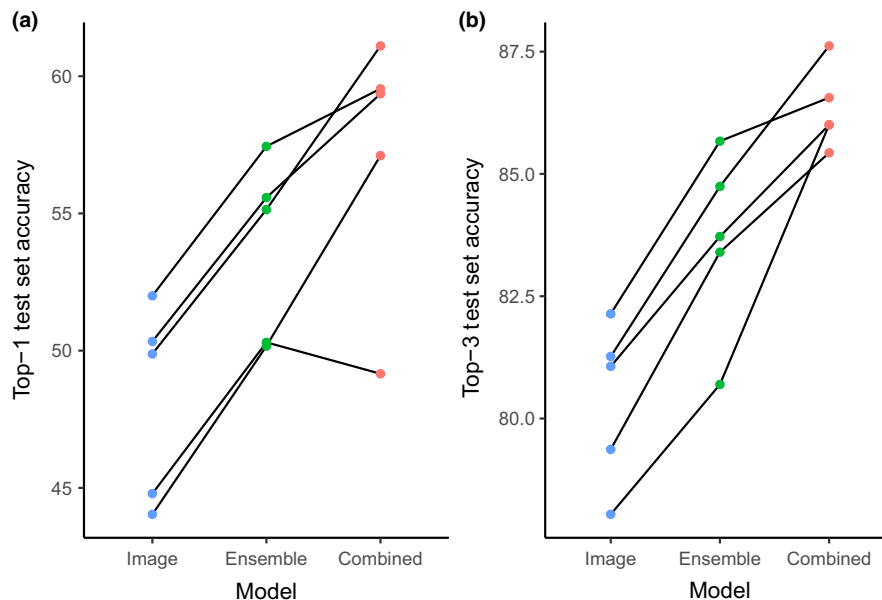


FIGURE 4 Consistent improvement in top-1 (a) and top-3 (b) accuracy from image-only models to models with the incorporation of metadata. An image-only model can be improved by ensembling with a metadata model, but further improvements can be gained from fitting combined multi-input models. Lines show 5 suites of models trained on a different train-validation-test randomizations. Mean improvement as model complexity is increased, and statistical significance determined from paired one-sided *t* tests were as follows (I = Image only, C = Combined, E = Ensemble): Top-1 I-E = +5.52 ($p < .0001$), E-C +3.53 ($p = .035$); I-C +9.05 ($p = .0019$); Top-3: I-C +3.23 ($p = .0001$), E-C +2.68 ($p = .011$), I-C +5.95 ($p = .0003$)

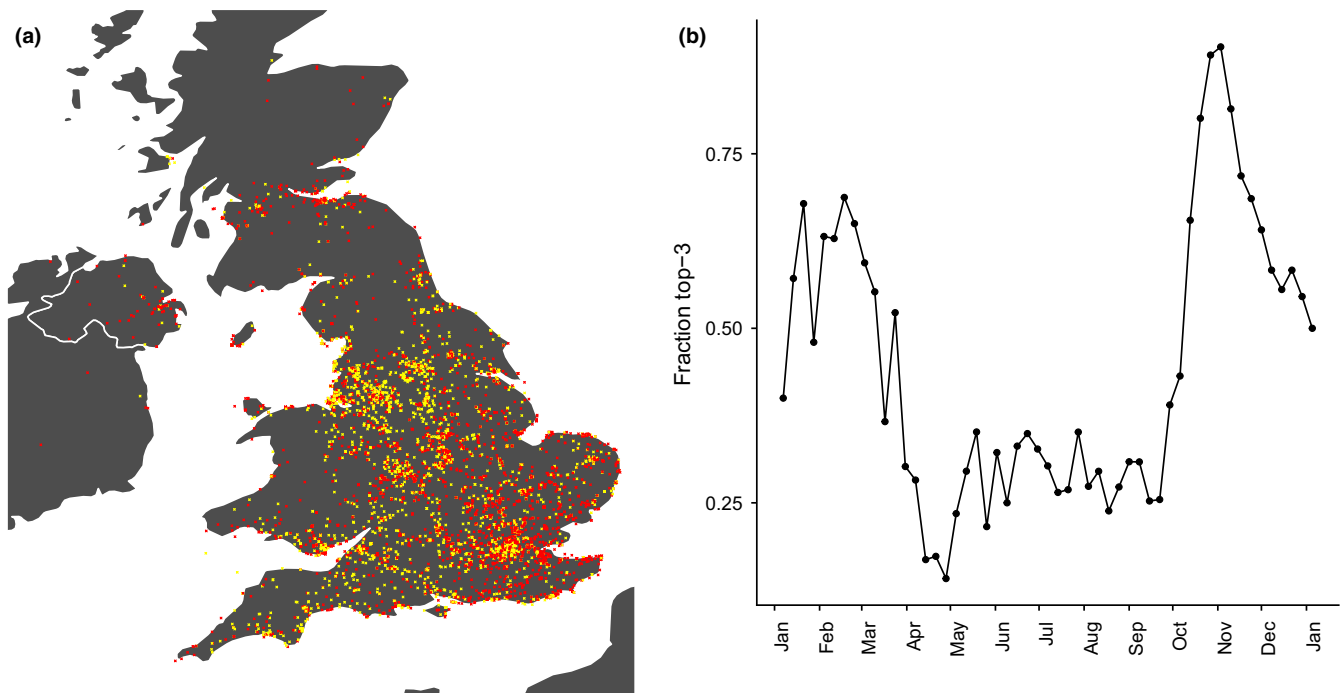


FIGURE 5 Distribution of records accurately (top-3) predicted solely from a derived metadata model. (a) Spatial distribution of accuracy, showing decreased accuracy in the south-east. Accurate predictions are shown in yellow, incorrect in red. (b) Weekly fraction of accurate metadata identifications through the year showing strong seasonal variation in accuracy with a particular peak in mid-autumn

appeared to be making more use of the weekly temperature data (2–10 weeks before the record), where randomization caused an 8.1 percentage-point decrease than the more proximate daily records for the preceding fortnight (5.4% decrease). The remaining metadata components had smaller influences on overall top-3 accuracy: randomising habitat data led to a 2.8% decrease while randomizing recorder experience led to a 3.1 percentage-point decrease.

These overall results are highly influenced by the dominant species (particularly the Harlequin ladybird) in the test set, masking variation in decline in accuracy on a per-species level (Table S1). The apparent importance of each metadata component appears to align with ecological expectations. The five species with greatest decline in accuracy when habitat is randomized are all considered habitat specialists (Roy & Brown, 2018): *Coccinella undecimpunctata* (dunes), *Anatis ocellata* (conifers), *Tytthaspis sedecimpunctata* (grassland and dunes), *Subcoccinella vigintiquatuorimpunctata* (grassland), and *Aphidecta oblitterata* (conifers). Similarly, the randomization of location had the greatest effect on the localized species (Figure S1). The top three most affected were: *Aphidecta oblitterata* (frequently reported in Scotland), *Scymnus interruptus* (South-East England) and *Coccinella undecimpunctata* (coastal). By contrast, *Coccinella septempunctata*, a widespread and generalist species was poorly identified by the metadata model and showed a minimal response to randomization. The species affected most by the randomization of temperature was *Propylea quatuordecimpunctata*, with the common name of the 'dormouse' ladybird (Roy & Brown, 2018, p. 112) because of its known late emergence.

The randomization of recorder experience had the greatest impact on *Scymnus interruptus* (a 33.6 percentage-point decrease). This

was the only 'inconspicuous' ladybird in our dataset, which inexperienced recorders may not even realize is a ladybird (see Figure 2g, right column). There was also a 5.5 percentage-point decrease in the identification of Harlequin ladybirds when recorder experience was randomized. Novice recorders are notably more likely to record Harlequin ladybirds than more experienced recorders. The first record submitted by a new recorder is a Harlequin ladybird 57.4% of the time, which rapidly declines to 38% by the 10th.

4 | DISCUSSION

The use of metadata within computer vision models considerably improves their reliability for species identification. This exciting finding has implications for biological recording, demonstrating the potential to use innovative approaches to assist in processing large occurrence datasets accrued through mass participation citizen science. Basic primary metadata is straightforward to incorporate within machine learning models and, since this information is already collected alongside the biological records, can be widely adopted.

4.1 | Interpretation of results

The notable gain in accuracy of the combined multi-input model compared to the ensemble model is consistent with the model learning to interpret the image based on the metadata. This is evidence that metadata can provide further gains beyond simply filtering the potential species list (Wittich et al., 2018). Further

evidence can be derived from the change in the relative confidence assigned to the true classification when metadata is incorporated (Appendix 3).

While it is not possible to determine exactly what interpretations the artificial intelligence is making, we can discern plausible scenarios. In autumn, ladybirds select suitable overwintering sites and enter dormancy through the adverse months (Roy & Brown, 2018). Each species exhibits a specific preference in overwintering habitat. Harlequin ladybirds favour buildings, leading to a high proportion of submitted records from inside homes of Harlequin ladybirds in the autumn as they move inside to overwinter (Roy et al., 2016). Submitted images of ladybirds exhibiting this behaviour are often poor-quality showing ladybirds at a distance nestled in crevices (Figure 2). The high accuracy of the metadata model during autumn suggests it has learnt (as expert human verifiers have) that a poor-quality image with a pale background during the autumn is very likely a Harlequin ladybird.

Our results likely represent a lower bound on the potential improvements that can be leveraged from metadata for identifying challenging species. Although British ladybirds have distinct ranges, activity periods and habitat (Comont et al., 2012; Roy & Brown, 2018) many are relatively cosmopolitan and can be observed as adults for large parts of the year. Classification models where focal species are more localized in time, space or habitat, or alternatively if the domain of the model is larger (for example North America, Berg et al., 2014), may expect to see larger gains through including metadata.

Determining how deep learning models make decisions is complex (Goodfellow et al., 2016). Multiple interwoven contributing factors combine to produce a result, much akin to human decisions. The nature of metadata means much of the gain likely comes from ruling species out rather than positively identifying them, which makes the interpretation of 'accuracy' metrics even more challenging. Our randomization analysis to determine the features used by the metadata model can only be a rough guide to the basis of decisions. The randomization process will represent the pre-existing imbalance of our dataset and will produce illogical combinations of metadata, such as hot temperatures during the winter, or coastal habitat within inland areas. Nonetheless, it does show evidence that the model operates along similar lines to expert identifiers. Where certain aspects of information are lost, this translated into inaccuracies in species for which that information is relevant. This is aligned with the results of Miao et al. (2018) who found that their image recognition tool for savanna mammals also used similar features to humans to identify species. Equally, for widespread and generalist species, metadata is not able to contribute to the accuracy. For instance, the identification of Seven-spot ladybird is essentially unchanged by the inclusion of metadata.

In theory, given enough records, a deep-learning model would be able to infer the information content of the cross-referenced database based only on primary metadata. For example, a neural network could learn to identify a set of location coordinates with a high likelihood of a given species, without knowing that those

coordinates contained favoured habitat, simply because the species is frequently recorded at these locations in the training dataset. In this respect, the inclusion of derived metadata could be considered a feature extractor technique that interprets the primary metadata, rather than providing additional information. In practice, the level of data required to internally reconstruct sufficient mapping purely from primary metadata would be very high, particularly when the features are very high resolution (Tang et al., 2015). A core challenge for automated species identification is the long tail of species for which there are very sparse records (Van Horn et al., 2017), for which the advantage of including derived metadata is likely to be considerably larger than for frequently recorded species.

4.2 | Further improvements to model

The design and training of deep learning models is an art rather than an exact science (Chollet & Allaire, 2018). There are likely to be opportunities for improvement in overall accuracy for each of our models. Our image-only accuracy levels (48.2%) were below that attained on other ecological datasets, though citizen scientists' images of ladybirds have been previously identified as posing a particular challenge for computer vision systems (Van Horn et al., 2017). For example, 67% accuracy was established as a baseline on the diverse iNaturalist image competition dataset (Van Horn et al., 2017), while competition winners were able to reach 74%.

Practically, incorporating metadata into neural networks need not introduce considerably more effort. Metadata is substantially simpler to process than image data and did not appear to add significantly to the training time. Compared to the very deep convolutional networks needed to interpret images, metadata can be processed with a small number of densely connected layers. Our tests with much larger or deeper networks did not lead to further gains. The number of parameters in our metadata models was several orders of magnitude smaller than the image model and could be trained in a matter of seconds per epoch. However, there are small additional design overheads in constructing a multi-input neural network compared to an image-only approach. There now exist user-friendly 'automatic learning' software that can generate a computer vision model given only a set of labelled images. In contrast, currently available support for multi-input models is comparatively lacking and requires direct specification of the model architecture as well as data manipulation pipelines to combine disparate information sources. Fortunately, tools such as the KERAS R package (Allaire & Chollet, 2019) provide straightforward frameworks for multi-input models that are well within the reach of ecologists without a formal computational science background. We have also shared our code (Supporting Information) to help others make use of this methodology.

We have demonstrated the improvement gained through the use of metadata. Further improvements in accuracy could likely be made through instigating test-time augmentation where multiple crops or rotations of an image are presented to the classifier,

ensembling multiple models, and increasing the size of the dataset through supplementary images and historical records (Chollet & Allaire, 2018). Our approach to augmenting metadata (adding Gaussian noise to each element) was relatively basic and more targeted approaches to generating additional synthetic training data (Chawla, Bowyer, Hall, & Kegelmeyer, 2002) could lead to better results.

The overall accuracy of a species classifier can be considerably enhanced by incorporating a prior likelihood of each species' relative frequency. Approaches that allow the model to directly learn the relative frequencies of the species could attain even higher overall accuracy. However, in contrast to improvements discussed in the previous paragraph, this would significantly reduce the accuracy for rarely observed species. A model that only learnt to accurately distinguish between Harlequin and Seven-spot ladybirds (that constitute the majority of records) could attain an accuracy of 70%, but this would be of limited applied use.

The challenge of species identification has in the past attracted computer scientists who can view species identification as an interesting example of large real-world labelled datasets (Weinstein, 2018). Open competitions such as the annual iNaturalist (Van Horn et al., 2017) and LifeCLEF competitions (Goëau, Bonnet, & Joly, 2017) have spurred considerable improvements in identification accuracy. Including metadata in these datasets (such as the PlantCLEF 2019 competition) could lead to considerable improvements. However, any release of metadata must consider the geoprivacy of citizen scientists and potential risk to endangered species. Due consideration of the appropriate resolution of location data, and the identifiability of individuals in any data publicly released is essential.

4.3 | Transferability of models including metadata

The inclusion of metadata in an automatic identification tool will influence its transferability to new contexts. With all machine learning approaches, any automatic identification process is only as good as the extent and scope of the training data used. A model that has been trained on the location of UK records would need to be re-trained for use in continental Europe, whereas an image-only model could be expected to be at least somewhat useful in both contexts. As such, a model trained on derived metadata such as habitat types or local weather may be more transferable than one trained on coordinates and specific dates. Understanding the domain a model will be applied to is essential. Transferability will be critical for expanding from well-studied areas (such as UK), to understudied areas where there is great potential for citizen science to fill gaps in knowledge (Pocock et al., 2018).

Transferability of models can be a challenge even within a region since records generated through unstructured broad-based citizen science are distinctive from those generated by committed amateur recorders, structured citizen science projects or professional surveys (Boakes et al., 2016). Submitted records are

the result of interactions between human behaviour and species ecology (Boakes et al., 2016). Highly visited sites may show an over-abundance of common species that are new to citizen scientists with relatively limited experience. In our dataset, uploaded records of ladybirds correlate strongly with the first appearance of species and news reports of invasive species (M. Logie & T. A. August, unpublished data). In comparison to ecological data, the inclusion of observer behaviour needs to be treated with particular care. While 'ecological' factors could be expected to transfer well between datasets, observer behaviour is likely to be considerably less transferable. Nevertheless, when working with citizen science data, including observer behaviour can provide additional information (Johnston, Fink, Hochachka, & Kelling, 2018). In our dataset, we could gain additional information at either end of the reviewer experience spectrum—novice recorders were much more likely to record Harlequin ladybirds. There is also potential for more detailed metrics, such as observer range, frequency or previous identification accuracy, could further improve model accuracy.

Our choice of what contextual data to include was guided by our knowledge of variables that are likely to influence ladybirds in the British Isles. For more taxonomically diverse tools, it would be beneficial to use a wider range of derived metadata variables. This could include more diverse weather information, climate maps, and topography. We did not include species range maps (Roy, Brown, Frost, & Poland, 2011) in this study since most (>90%) records came from areas within the range of 15 out of the 18 focal species considered in this study. Binary species range maps cannot account for the relative frequency of species across a region, but this can be learnt by a deep learning network provided with location data of records. Although range maps could be informative within models with a wide spatial scope or for highly localized species, they are comparatively verbose to encode for in deep learning networks. When using a model to identify large numbers of species, the intersection or otherwise of a record with each species range map may need to be encoded in a separate variable. This greatly increases the length of the metadata vector associated with each record and it could become challenging for models to identify relevant information. Although deep learning networks have the potential to effectively ignore data that is not relevant, there is the potential to slow the fitting procedure if too much irrelevant information is presented. Where accurate species range map data are available (and may impart additional information beyond that contained in the training set of records), an approach that combines machine learning with a range-map-based shortlist may be the most useful (Wittich et al., 2018).

5 | CONCLUSIONS

Identification of insects poses a considerable challenge for computer vision (Martineau et al., 2017). Insect diversity is extraordinarily large – as an example, there are over 6,000 ladybird species worldwide (Roy & Brown, 2018), most of which do not have

accessible labelled images. For difficult challenges, such as species identification in the field, the optimal solutions will involve humans and artificial intelligence working in tandem (Trouille, Lintott, & Fortson, 2019). Our results demonstrate the potential for considerable improvement in the accuracy of automatic identification when incorporating contextualization information directly within the model. This is also likely to apply to passive acoustic monitoring tools (Gibb, Browning, Glover-Kapfer, & Jones, 2019) too. Researchers building automatic identification methods will benefit from training models to place images in context, just as a human naturalist would, to best unlock the potential of artificial intelligence in ecology.

ACKNOWLEDGEMENTS

Our thanks to Mark Logie for assistance accessing the iRecord database, Colin Harrower for species range maps, and to the UK Ladybird Survey volunteer recorders who generated the dataset the work is based upon. This work was supported by the Natural Environment Research Council award number NE/R016429/1 as part of the UK-SCAPE programme delivering National Capability. J.C.D.T. was funded for this project through an NERC Innovation placement linked to the Oxford Environmental Doctoral Training Partnership (NE/L002612/1). We thank two reviewers for their constructive comments.



AUTHORS' CONTRIBUTIONS

J.C.D.T. built the models and analysed the results, based on an initial idea and design of T.A.A. and through discussions with H.E.R. and T.A.A. J.C.D.T. wrote the first manuscript draft and all authors contributed critically to revisions.

DATA AVAILABILITY STATEMENT

R code used to build models and analyse results is available at https://github.com/jcdterry/LadybirdID_Public (<https://doi.org/10.5281/zenodo.3530383>) (Terry, 2019) and summarized in Supporting Information 2. Images, user IDs and location data used in this paper are not publically archived due to image licensing and data protection constraints. Data can be accessed from the Biological Records Centre indicia database by searching for records in the family Coccinellidae, or with a common name that included 'ladybird', that were not marked as rejected or dubious. For access to the indicia database contact brc@ceh.ac.uk.

ORCID

J. Christopher D. Terry  <https://orcid.org/0000-0002-0626-9938>
 Helen E. Roy  <https://orcid.org/0000-0001-6050-679X>
 Tom A. August  <https://orcid.org/0000-0003-1116-3385>

REFERENCES

- Allaire, J., & Chollet, F. (2019). keras: R Interface to 'Keras'. Retrieved from <https://cran.r-project.org/package=keras>
- August, T., Harvey, M., Lightfoot, P., Kilbey, D., Papadopoulos, T., & Jepson, P. (2015). Emerging technologies for biological recording. *Biological Journal of the Linnean Society*, 115(3), 731–749. <https://doi.org/10.1111/bij.12534>
- Berg, T., Liu, J., Lee, S. W., Alexander, M. L., Jacobs, D. W., & Belhumeur, P. N. (2014). Birdsnap: Large-scale fine-grained visual categorization of birds. *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 2019–2026. <https://doi.org/10.1109/CVPR.2014.259>
- Boakes, E. H., Gliozzo, G., Seymour, V., Harvey, M., Smith, C., Roy, D. B., & Haklay, M. (2016). Patterns of contribution to citizen science biodiversity projects increase understanding of volunteers' recording behaviour. *Scientific Reports*, 6(1), 33051. <https://doi.org/10.1038/srep33051>
- Buda, M., Maki, A., & Mazurowski, M. A. (2018). A systematic study of the class imbalance problem in convolutional neural networks. *Neural Networks*, 106, 249–259. <https://doi.org/10.1016/j.neunet.2018.07.011>
- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321–357. <https://doi.org/10.1613/jair.953>
- Chollet, F., & Allaire, J. (2018). *Deep Learning in R*. Shelter Island, NY: Manning.
- Christin, S., Hervet, É., & Lecomte, N. (2019). Applications for deep learning in ecology. *Methods in Ecology and Evolution*, 10(10), 1632–1644. <https://doi.org/10.1111/2041-210X.13256>
- Comot, R. F., Roy, H. E., Lewis, O. T., Harrington, R., Shortall, C. R., & Purse, B. V. (2012). Using biological traits to explain ladybird distribution patterns. *Journal of Biogeography*, 39(10), 1772–1781. <https://doi.org/10.1111/j.1365-2699.2012.02734.x>
- Creed, E. R. (1966). Geographic variation in the two-spot Ladybird in England and Wales. *Heredity*, 21(1), 57–72. <https://doi.org/10.1038/hdy.1966.4>
- Cui, Y., Song, Y., Sun, C., Howard, A., & Belongie, S. (2018). Large scale fine-grained categorization and domain-specific transfer learning. *arXiv preprint*. <https://doi.org/10.1109/CVPR.2018.00432>
- Duff, A. G. (2018). *Checklist of beetles of the British Isles* (3rd ed.). Iver: Pemberley Books.
- Ellen, J. S., Graff, C. A., & Ohman, M. D. (2019). Improving plankton image classification using context metadata. *Limnology and Oceanography: Methods*, 17(8), 439–461. <https://doi.org/10.1002/lom3.10324>
- Gardiner, M. M., Allee, L. L., Brown, P. M., Losey, J. E., Roy, H. E., & Smyth, R. R. (2012). Lessons from lady beetles: Accuracy of monitoring data from US and UK citizen-science programs. *Frontiers in Ecology and the Environment*, 10(9), 471–476. <https://doi.org/10.1890/110185>
- Gaston, K. J., & O'Neill, M. A. (2004). Automated species identification: Why not? *Philosophical Transactions of the Royal Society B: Biological Sciences*, 359(1444), 655–667. <https://doi.org/10.1098/rstb.2003.1442>
- Gibb, R., Browning, E., Glover-Kapfer, P., & Jones, K. E. (2019). Emerging opportunities and challenges for passive acoustics in ecological assessment and monitoring. *Methods in Ecology and Evolution*, 10(2), 169–185. <https://doi.org/10.1111/2041-210X.13101>
- Goëau, H., Bonnet, P., & Joly, A. (2017). Plant identification based on noisy web data: The amazing performance of deep learning. In C. Linda, F. Nicola, G. Lorraine, & M. Thomas (Eds.), *CEUR Workshop Proceedings*. Retrieved from <http://ceur-ws.org/Vol-1886/>
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. Cambridge, MA: MIT Press.
- IPBES. (2019). *Global assessment report on biodiversity and ecosystem services of the Intergovernmental Science-Policy Platform on Biodiversity and Ecosystem Services* (E. S. Brondizio, J. Settele, S. Díaz, & H. T. Ngo, Eds.). Bonn, Germany: IPBES Secretariat.
- Isaac, N. J. B., van Strien, A. J., August, T. A., de Zeeuw, M. P., & Roy, D. B. (2014). Statistics for citizen science: Extracting signals of change from noisy ecological data. *Methods in Ecology and Evolution*, 5(10), 1052–1060. <https://doi.org/10.1111/2041-210X.12254>
- Johnston, A., Fink, D., Hochachka, W. M., & Kelling, S. (2018). Estimates of observer expertise improve species distributions from citizen

- science data. *Methods in Ecology and Evolution*, 9(1), 88–97. <https://doi.org/10.1111/2041-210X.12838>
- Jouveau, S., Delaunay, M., Vignes-Lebbe, R., & Nattier, R. (2018). A multi-access identification key based on colour patterns in ladybirds (Coleoptera, Coccinellidae). *ZooKeys*, 758, 55–73. <https://doi.org/10.3897/zookeys.758.22171>
- Kingma, D. P., & Ba, J. (2014). ADAM: A method for stochastic optimization. *ArXiv Preprint*, 1–15. Retrieved from <http://arxiv.org/abs/1412.6980>
- Laina, I., Rupprecht, C., Belagiannis, V., Tombari, F., & Navab, N. (2016). Deeper depth prediction with fully convolutional residual networks. *Proceedings – 2016 4th International Conference on 3D Vision, 3DV 2016* (pp. 239–248). <https://doi.org/10.1109/3DV.2016.32>
- Mac Aodha, O., Cole, E., & Perona, P. (2019). Presence-only geographical priors for fine-grained image classification. Retrieved from <http://arxiv.org/abs/1906.05272>
- Marques, A. C. R., Raimundo, M. M., Cavalheiro, E. M. B., Salles, L. F. P., Lyra, C., & Von Zuben, F. J. (2018). Ant genera identification using an ensemble of convolutional neural networks. *PLoS ONE*, 13(1), 1–13. <https://doi.org/10.1371/journal.pone.0192011>
- Martineau, M., Conte, D., Raveaux, R., Arnault, I., Munier, D., & Venturini, G. (2017). A survey on image-based insect classification. *Pattern Recognition*, 65, 273–284. <https://doi.org/10.1016/j.patcog.2016.12.020>
- Martineau, M., Raveaux, R., Chatelain, C., Conte, D., & Venturini, G. (2018). Effective training of convolutional neural networks for insect image recognition. In J. Blanc-Talon, D. Helbert, W. Philips, D. Popescu, & P. Scheunders (Eds.), *Advanced concepts for intelligent vision systems. ACIVS 2018. Lecture Notes in Computer Science*. (Vol. 11182, pp. 426–437). Cham: Springer. https://doi.org/10.1007/978-3-030-01449-0_36
- Met Office. (2012). Met Office Integrated Data Archive System (MIDAS) land and marine surface stations data (1853–current). NCAS British Atmospheric Data Centre 2019. Retrieved from <http://catalogue.ceda.ac.uk/uuid/220a65615218d5c9cc9e4785a3234bd0>
- Miao, Z., Gaynor, K. M., Wang, J., Liu, Z., Muellerklein, O., Norouzzadeh, M. S., ... Getz, W. M. (2018). A comparison of visual features used by humans and machines to classify wildlife. *BioRxiv*, 450189. <https://doi.org/10.1101/450189>
- Minetto, R., Pamplona Segundo, M., & Sarkar, S. (2019). Hydra: An ensemble of convolutional neural networks for geospatial land classification. *IEEE Transactions on Geoscience and Remote Sensing*, 57(9), 6530–6541. <https://doi.org/10.1109/TGRS.2019.2906883>
- Pocock, M. J. O., Chandler, M., Bonney, R., Thornhill, I., Albin, A., August, T., Danielsen, F. (2018). A vision for global biodiversity monitoring with citizen science. In D. A. Bohan, A. J. Dumbrell, G. Woodward, & M. Jackson (Eds.), *Advances in ecological research* (pp. 169–223). London, UK: Academic Press. <https://doi.org/10.1016/bs.aecr.2018.06.003>
- Pocock, M. J. O., Roy, H. E., Preston, C. D., & Roy, D. B. (2015). The Biological Records Centre: A pioneer of citizen science. *Biological Journal of the Linnean Society*, 115(3), 475–493. <https://doi.org/10.1111/bj.12548>
- Rowland, C. S., Morton, R. D., Carrasco, L., McShane, G., O'Neil, A. W., & Wood, C. M. (2017). Land Cover Map 2015 (1 km percentage target class, GB). NERC Environmental Information Data Centre. <https://doi.org/10.5285/505d1e0c-ab60-4a60-b448-68c5bbae403e>
- Roy, H. E., & Brown, P. (2018). *Field guide to the ladybirds of Great Britain and Ireland*. London, UK: Bloomsbury.
- Roy, H. E., Brown, P. M. J., Adriaens, T., Berkvens, N., Borges, I., Clusella-Trullas, S., ... Zhao, Z. (2016). The harlequin ladybird, *Harmonia axyridis*: Global perspectives on invasion history and ecology. *Biological Invasions*, 18(4), 997–1044. <https://doi.org/10.1007/s10530-016-1077-6>
- Roy, H. E., Brown, P. M., Frost, R., & Poland, R. L. (2011). *Ladybirds (Coccinellidae) of Britain and Ireland*. Shrewsbury: Field Studies Council.
- Rzanny, M., Seeland, M., Wäldchen, J., & Mäder, P. (2017). Acquiring and preprocessing leaf images for automated plant identification: Understanding the tradeoff between effort and information gain. *Plant Methods*, 13(1), 1–11. <https://doi.org/10.1186/s13007-017-0245-8>
- Silvertown, J. (2009). A new dawn for citizen science. *Trends in Ecology and Evolution*, 24(9), 467–471. <https://doi.org/10.1016/j.tree.2009.03.017>
- Swanson, A., Kosmala, M., Lintott, C., Simpson, R., Smith, A., & Packer, C. (2015). Snapshot Serengeti, high-frequency annotated camera trap images of 40 mammalian species in an African savanna. *Scientific Data*, 2, 150026. <https://doi.org/10.1038/sdata.2015.26>
- Szegedy, C., Ioffe, S., Vanhoucke, V., & Alemi, A. (2016). Inception-v4, Inception-ResNet and the impact of residual connections on learning. *ArXiv*. <https://doi.org/10.1016/j.patrec.2014.01.008>
- Tang, K., Paluri, M., Fei-Fei, L., Fergus, R., & Bourdev, L. (2015). Improving image classification with location context. *Proceedings of the IEEE International Conference on Computer Vision, 2015 Inter* (pp. 1008–1016). <https://doi.org/10.1109/ICCV.2015.121>
- Terry, C. (2019). jcdterry/LadybirdID_Public: MEE Publication (Version vol 1.0). Zenodo. <http://doi.org/10.5281/zenodo.3530383>
- Torney, C. J., Lloyd-Jones, D. J., Chevallier, M., Moyer, D. C., Maliti, H. T., Mwita, M., ... Hopcraft, G. C. (2019). A comparison of deep learning and citizen science techniques for counting wildlife in aerial survey images. *Methods in Ecology and Evolution*, 10(6), 779–787. <https://doi.org/10.1111/2041-210X.13165>
- Trouille, L., Lintott, C. J., & Fortson, L. F. (2019). Citizen science frontiers: Efficiency, engagement, and serendipitous discovery with human-machine systems. *Proceedings of the National Academy of Sciences*, 116(6), 1902–1909. <https://doi.org/10.1073/pnas.1807190116>
- Van Horn, G., Mac Aodha, O., Song, Y., Cui, Y., Sun, C., Shepard, A., ... Belongie, S. (2017). The iNaturalist species classification and detection dataset. *ArXiv*. Retrieved from <http://arxiv.org/abs/1707.06642>
- Wäldchen, J., & Mäder, P. (2018). Machine learning for image based species identification. *Methods in Ecology and Evolution*, 9(11), 2216–2225. <https://doi.org/10.1111/2041-210X.13075>
- Weinstein, B. G. (2018). A computer vision for animal ecology. *Journal of Animal Ecology*, 87(3), 533–545. <https://doi.org/10.1111/1365-2656.12780>
- Willi, M., Pitman, R. T., Cardoso, A. W., Locke, C., Swanson, A., Boyer, A., ... Fortson, L. (2019). Identifying animal species in camera trap images using deep learning and citizen science. *Methods in Ecology and Evolution*, 10(1), 80–91. <https://doi.org/10.1111/2041-210X.13099>
- Wittich, H. C., Seeland, M., Wäldchen, J., Rzanny, M., & Mäder, P. (2018). Recommending plant taxa for supporting on-site species identification. *BMC Bioinformatics*, 19(1), 1–17. <https://doi.org/10.1186/s12859-018-2201-7>

SUPPORTING INFORMATION

Additional supporting information may be found online in the Supporting Information section.

How to cite this article: Terry JCD, Roy HE, August TA.

Thinking like a naturalist: Enhancing computer vision of citizen science images by harnessing contextual data. *Methods Ecol Evol*. 2020;11:303–315. <https://doi.org/10.1111/2041-210X.13335>